

ORBIT Capability Matrix

Overview

ORBIT (Open Retrieval-Based Inference Toolkit) is an open-source AI gateway, retrieval-augmented generation (RAG) engine, and agent-protocol host. It can be deployed on-premises, in a private cloud, or fully offline.

Models, data sources, adapters, and security settings are defined as YAML files under `config/`, rather than in server code. This allows configuration to be version-controlled and reviewed in a source repository like any other code artifact, diffed between commits, and promoted across development, staging, and production environments through the organization's existing CI/CD and change-management process — supporting typical DevSecOps practices for multi-environment deployment.

This document describes ORBIT's capabilities as implemented in its configuration files (`config/*.yaml`) and documentation (`docs/`), and compares them against three reference categories:

- **LiteLLM** — an LLM API proxy and router
- **Open WebUI** — a chat user interface for local/cloud LLMs
- **Commercial managed AI platforms** (e.g., AWS Bedrock, Azure AI, Google Vertex AI, IBM watsonx)

Each row in the comparison tables is intended to describe what each platform does, not to rank them. Where a claim about a comparison platform could not be verified against that platform's own documentation or pricing page, it has been noted as such or omitted.

1. Comparison Matrix

Dimension	ORBIT	LiteLLM	Open WebUI	Commercial Managed Platforms (Bedrock/Azure/Vertex)
Primary function	AI gateway + multi-source RAG + agent protocol host	LLM API proxy and router	Chat user interface for local/cloud LLMs	Managed cloud AI service
Inference backends	41 provider/runtime configurations (cloud APIs, local GGUF, vLLM, TensorRT-LLM, BitNet, AirLLM, TEE) — see note in §2.1 on default-disabled backends	Broad set of cloud/proxy API routes	Ollama-native; also reaches any OpenAI-compatible endpoint (OpenAI, or a self-hosted llama.cpp/vLLM/LM Studio server)	Vendor-operated model catalog
Data source integration	SQL (9 dialects), vector DBs (6 integrated), NoSQL, REST/GraphQL, web search (provider-native for Gemini/OpenAI/xAI, plus 7 external search backends usable with any LLM)	No built-in data-source connectors; typically paired with an external RAG pipeline	Web search plugin and document-based vector RAG	Cloud-native connectors (S3/Blob/GCS) and managed vector search services
Document & media processing	Three document-parsing engines (Docling, MarkItDown, LLM/vision OCR); image, video, STT/TTS generation	Not included	Text extraction from uploaded files (e.g., PyPDF)	Cloud OCR/vision services (e.g., Textextract, Form Recognizer)
Agent protocols	MCP (server and client host), Google A2A, RabbitMQ-based async messaging	Function/tool-call proxying	Tool-calling UI plugins	Vendor-specific agent frameworks
Conversational UX features	Cross-adapter skill invocation with optional automatic intent routing, query autocomplete from adapter/skill examples, and per-model dynamic conversation-history budgeting (see §2.9)	Not in scope — LiteLLM proxies inference calls and does not manage conversation state or UX	Chat UI includes model-side features (e.g., prompt templates, RAG toggles) but does not implement adapter-swapping skills, intent-based auto-routing, or dynamic per-model history budgeting	Varies by vendor SDK/console; not a standard cross-platform feature
Security & identity	6-role RBAC (11 permissions), OIDC/SSO (Entra ID, Auth0), OS keyring integration, AES-256 file encryption — included in the open-source distribution	Key-based limits and proxy-level auth in the open-source tier; SSO/SCIM, OIDC/JWT auth, and audit logs are listed as part of the paid Enterprise tier per LiteLLM's published pricing (see §3.1)	User accounts, OAuth	Cloud IAM (AWS IAM, Azure RBAC, etc.)

Dimension	ORBIT	LiteLLM	Open WebUI	Commercial Managed Platforms (<i>Bedrock/Azure/Vertex</i>)
Safety & moderation	Configurable moderation backends (OpenAI, Anthropic, Llama-Guard3, Shieldstral), plus two PII detectors: a local model (<code>privacy-filter</code>) and a Presidio analyzer integration	Integrates with external moderation/PII services (e.g., Presidio)	Basic moderation options	Managed cloud guardrail services
Hardware acceleration	CUDA, MPS, vLLM, TensorRT-LLM (FP8/INT8), SGLang, BitNet	Routes to external inference servers; does not run hardware-accelerated inference itself	Dependent on the Ollama backend it connects to	Managed cloud GPU infrastructure
Deployment model	Docker, Kubernetes, bare-metal, air-gapped, async worker processes	Docker, Kubernetes, serverless	Docker, desktop, serverless	Managed cloud SaaS/PaaS
License	Apache 2.0 (single tier, no gated features)	Open-source core under a permissive license, with a separate paid Enterprise tier for the features listed in §3.1 (per litellm.ai/pricing)	BSD-3-Clause through v0.6.5; v0.6.6+ adds a branding-protection clause (not OSI-approved as "open source"), with a separate enterprise license for white-labeling — see §4.2	Proprietary cloud service

2. Capability Detail

2.1 LLM Gateway & Inference

ORBIT defines **41 inference provider/runtime configurations** in `config/inference.yaml` . Most are disabled by default and require an API credential or a locally running runtime (e.g., an Ollama daemon or a vLLM server) to be usable — the count reflects configured integrations, not backends active in a given deployment.

Capability	Details	Source
Supported backends (41)	Ollama, vLLM, SGLang, TensorRT-LLM, Shimmy, llama.cpp, Gemini, Groq, DeepSeek, Vertex AI, AWS Bedrock, Azure OpenAI, OpenAI, Mistral, Anthropic, Together AI, xAI Grok, Hugging Face Transformers, AirLLM, OpenRouter, Cohere, IBM watsonx,	<code>config/inference.yaml</code>

Capability	Details	Source
entries)	Perplexity, Fireworks AI, Replicate, NVIDIA AI Catalog, BitNet, ZAI, Cerebras, DeepInfra, LM Studio, Moonshot AI, MiniMax, NEAR AI Cloud (TEE), Nebius, Venice, Scaleway, Sakana Fugu	
Local hardware acceleration	TensorRT-LLM (FP8, INT8, AWQ 4-bit), vLLM (tensor/pipeline parallelism), SGLang, AirLLM (disk-based layer streaming)	config/inference.yaml
Quantized/low-bit runtimes	BitNet (<code>i2_s</code> , <code>t11</code> kernels), llama.cpp GGUF	config/inference.yaml , config/llama_cpp.yaml , docs/bitnet-setup.md
Confidential computing	NEAR AI Cloud TEE integration	config/inference.yaml
Fault tolerance	Circuit breaker with exponential backoff, jitter, recovery timeout, and probe caps (<code>failure_threshold: 5</code> , <code>recovery_timeout: 30s</code>); model fallback chaining	config/config.yaml (<code>fault_tolerance</code>)
Per-adapter provider overrides	Adapters can override <code>inference_provider</code> , <code>embedding_provider</code> , <code>reranker_provider</code> , and <code>model</code>	config/adapters.yaml

2.2 Data Sources & Retrieval

Capability	Details	Source
SQL databases (9 dialects)	PostgreSQL, MySQL, MariaDB, SQLite, Supabase, Oracle, SQL Server, DuckDB, AWS Athena	config/datasources.yaml
Vector databases (6 integrated as retrievers)	Chroma, Qdrant, Milvus, Pinecone, Elasticsearch, Redis Vector. Weaviate appears in <code>stores.yaml</code> for store-lifecycle management, not as an active retriever backend.	config/datasources.yaml , config/stores.yaml
NoSQL databases	MongoDB (Atlas or self-hosted), Cassandra	config/datasources.yaml
Web/API sources	REST, GraphQL	—
Web search — provider-native	Delegates search and synthesis to the LLM provider's own built-in tool in a single call. Supported with Gemini (<code>google_search</code> grounding), OpenAI (Responses API <code>web_search</code>), and xAI (<code>web_search</code>). Other providers are not supported for this mode.	config/adapters/web-search.yaml , docs/adapters/web-search.md

Capability	Details	Source
Web search — external providers	A dedicated <code>WebSearchStep</code> calls an external search API (DuckDuckGo, SearXNG, Brave, Serper, Tavily, Google PSE, Perplexity), then any configured LLM synthesizes the answer from the results. Decouples the search backend from the synthesizing model (e.g., DuckDuckGo results summarized by a local Ollama model).	<code>config/adapters/web-search-providers.yaml</code> , <code>docs/adapters/web-search.md</code>
Embedding providers (11)	Ollama, llama.cpp, Jina, OpenAI, Cohere, Mistral, Gemini, Voyage AI, OpenRouter, NVIDIA, SentenceTransformers	<code>config/embeddings.yaml</code>
Rerankers (7)	Ollama (BGE-Reranker v2), Cohere, Jina, OpenAI, Anthropic, Voyage AI, OpenRouter	<code>config/rerankers.yaml</code>
Hybrid retrieval scoring	Combines dense embedding similarity, lexical similarity (Jaro-Winkler, Levenshtein), and reranking	<code>config/config.yaml</code> (<code>composite_retrieval</code>)
SQL query safety guard	Generated SQL is validated before execution to confirm it is a single, read-only, row-capped statement, rejecting non-query commands	<code>CHANGELOG.md</code> (v2.15.3)
Intent template validation	Intent template libraries are validated against a formal schema at load time; templates can optionally require explicit approval before being served	<code>CHANGELOG.md</code> (v2.15.4)

Retrieval is implemented via **intent-based query generation**: natural-language input is mapped to SQL/NoSQL/API queries through configured templates, rather than a general-purpose text-to-SQL model call on every request.

2.3 Document Processing

Capability	Details	Source
File formats	PDF, DOCX, PPTX, XLSX, XLS, CSV, JSON, XML, HTML, TXT, EPUB, ZIP, images, audio	<code>config/config.yaml</code>
Processing engines (in priority order)	1. Docling (layout-aware parsing) 2. MarkItDown (document-to-markdown) 3. LLM/vision-based OCR (Mistral, Gemini, or vision models from OpenAI/Anthropic/Cohere/Ollama/vLLM/llama.cpp)	<code>config/config.yaml</code> (<code>files.processing.processor_priority</code>), <code>config/ocr.yaml</code>
Chunking strategies	Fixed-size, token-based, semantic (similarity-boundary detection), recursive, markdown-header	<code>config/config.yaml</code> (<code>files.default_chunking_strategy</code>)
Document security	Magika MIME-type verification, AES-256-GCM encryption at rest	<code>config/config.yaml</code> (<code>files.processing.magika</code> , <code>files.encryption</code>)

Capability	Details	Source
Storage backends	Local filesystem, AWS S3/MinIO, Azure Blob Storage, Google Cloud Storage	<code>config/config.yaml</code> (<code>files.storage_backend</code>)

2.4 Multimodal & Media Processing

Domain	Details	Source
Speech-to-text (12 configured)	Local Whisper (tiny-large-v3, CPU/CUDA), OpenAI API, Google STT, Gemini, xAI Grok STT, Cohere Transcribe, OpenRouter, Ollama; 4 of the 12 entries are provider placeholders without a default model set	<code>config/stt.yaml</code>
Text-to-speech (11 configured)	OpenAI TTS, Google TTS, Gemini TTS, ElevenLabs, Supertonic (local, 31 languages), Coqui TTS (local), vLLM Orpheus, Ollama (Piper/Kokoro), OpenRouter; 2 of the 11 entries are placeholders	<code>config/tts.yaml</code>
Audio content sanitization	Strips non-speech content (code blocks, tables) before TTS synthesis, with an optional spoken placeholder	<code>config/tts.yaml</code> (<code>tts.sanitize_content</code>)
Image generation (5 configured)	OpenAI (DALL-E 2/3, GPT Image), Gemini, xAI Grok Imagine, Ollama (Flux2/Z-Image Turbo), OpenRouter (Seedream)	<code>config/image.yaml</code>
Video generation (3 configured)	Gemini (Veo), xAI (Grok Imagine Video), OpenRouter (Seedance)	<code>config/video.yaml</code>
Vision/OCR analysis (8 configured)	OpenAI, Anthropic, Cohere, Ollama, vLLM, llama.cpp, Gemini, Mistral	<code>config/vision.yaml</code>
Real-time speech-to-speech	Native, interruptible voice conversations (not a separate STT-then-TTS pipeline) via OpenAI Realtime (<code>gpt-realtime</code>) and Gemini Live (<code>gemini-3.1-flash-live-preview</code>), with barge-in support so a user can interrupt the assistant mid-answer. Can be grounded in a configured retriever adapter (SQL, vector, or other data source) via a <code>grounding_adapter</code> reference, so spoken answers stay factual rather than relying only on the model's own knowledge — usable for live virtual-assistant scenarios.	<code>config/adapters/qa.yaml</code> (<code>qa-realtime-voice</code> , <code>qa-gemini-realtime-voice</code>), <code>docs/adapters/grounded-realtime-voice.md</code>

Per-adapter allowlists (`allowed_image_models` , `allowed_video_models` , `allowed_audio_models` , `allowed_search_providers`) let a client select, at request time, from a configured set of provider/model combinations for image, video, audio, and web-search generation, rather than being fixed to one adapter-configured default. Because OpenRouter is itself an aggregator exposing many third-party models through a single provider entry, an OpenRouter-backed adapter can expose a correspondingly larger set of selectable image/video/audio models than a single-vendor provider would, subject to how many are added to that adapter's allowlist. (`CHANGELOG.md` , v2.15.6)

2.5 Protocols

Protocol	Details	Source
REST / OpenAI-compatible	<code>/v1/chat</code> , <code>/v1/chat/completions</code> with SSE streaming	<code>server/routes /</code>
Model Context Protocol (MCP)	Operates as both an MCP server (<code>/mcp</code>) and an MCP client host. Pre-wired example client configs: <code>filesystem</code> , a sample REST API, and <code>github</code> . Other MCP servers (e.g., Slack, Jira, Sentry, M365) are documented as configuration templates, not pre-wired connections.	<code>config/mcp_clients.yaml</code> , <code>docs/mcp_protocol.md</code>
MCP agentic tool-calling	A bounded, multi-step tool-calling loop lets the model call one or more MCP tools, receive results, and continue reasoning across several rounds within a single request — implemented as ORBIT's own code path against each provider's native tool-calling API, with no external agent-orchestration library (e.g. LangChain, AutoGen, CrewAI) involved. Available either as an explicit skill the client opts into per request, or in an "opportunistic" mode where an ordinary conversational adapter decides per turn, with no client-side signal, whether a tool is needed. Supported with OpenAI, Anthropic, Gemini, xAI, Ollama, Ollama Cloud, llama.cpp (API mode), and vLLM (API mode).	<code>docs/adapters /mcp-agent.md</code>
Google A2A	Agent discovery (<code>/.well-known/agent.json</code>) and task execution (<code>/a2a</code> , JSON-RPC 2.0, streaming)	<code>docs/a2a-protocol.md</code>
Async messaging	RabbitMQ-based AMQP request/response surface for decoupled processing	<code>config/config.yaml</code> (<code>messaging</code>)
WebSockets	Bi-directional streaming for audio and telemetry	<code>/ws</code> , <code>/ws/metrics</code>

2.6 Security, Governance & Identity

Capability	Details	Source
Authentication	PBKDF2-SHA256 password hashing (600,000 iterations), 256-bit opaque bearer tokens, OS keyring integration	docs/authentication.md
SSO	OIDC/OAuth2 for Microsoft Entra ID and Auth0; admin panel SSO with PKCE	config/config.yaml (auth.providers)
RBAC	6 built-in roles (admin , operator , auditor , analyst , user-manager , user), 11 permissions	docs/rbac-architecture.md
API key controls	Scoped keys, daily/monthly quotas, rate limits, user/email restrictions, adapter-level binding	docs/api-keys.md
Identity blacklisting	Wildcard pattern blocking with session revocation	config/config.yaml (auth.blacklist)
Audit logging	Request and admin-action logs, storable in SQLite, Postgres, MongoDB, or Elasticsearch	config/config.yaml (internal_services.audit)
Secrets management	Resolves \${VAR} references from AWS Secrets Manager, Azure Key Vault, or GCP Secret Manager, falling back to .env	config/config.yaml (secrets_management)
Strict auth mode	Optional setting requiring a valid user bearer token in addition to an API key	config/config.yaml (auth.require_authenticated_user)

All items in this table are included in ORBIT's open-source (Apache 2.0) distribution; there is no separate paid tier.

2.7 Safety, Moderation & Privacy

Capability	Details	Source
Moderation backends	OpenAI Omni-Moderation, Anthropic moderator, Ollama Llama-Guard3, Shieldstral (via vLLM/llama.cpp)	config/moderators.yaml , config/guardrails.yaml
Local PII detection	Token-classification model (openai/privacy-filter , Apache 2.0) run on-premises; no outbound call required	config/moderators.yaml (privacy_filter)
PII detection via Presidio	Presidio analyzer called over its REST API (no added Python dependencies); ~100 configurable entity types, per-entity score thresholds, self-hosted	config/moderators.yaml (presidio), docs/security/pii-moderation.md

Capability	Details	Source
Network security controls	CORS, CSP, HSTS, X-Frame-Options, X-Content-Type-Options, rate limiting, progressive throttling	<code>config/config.yaml (security)</code>

2.8 Reliability & Performance

Capability	Details	Source
Multi-worker scaling	Uvicorn process scaling with separate thread pools per workload type	<code>config/config.yaml (performance)</code>
Caching	SQLite, Redis, or Memcached backends	<code>config/config.yaml (internal_services.cache)</code>
Response optimization	GZip compression, ETag caching	<code>config/config.yaml (performance.compression , performance.etag_caching)</code>
Cost/token tracking	Per-request token usage and estimated cost, by provider/model	<code>config/pricing.yaml</code> , <code>docs/token-usage-and-cost-tracking.md</code>
Observability	WebSocket telemetry, Prometheus metrics endpoint, log rotation	<code>config/config.yaml</code>

2.9 Conversational Behavior & Adapter Intelligence

Capability	Details	Source
Skills (cross-adapter capability invocation)	Any adapter can invoke another adapter's function mid-conversation via a <code>skill</code> request field (e.g., a retrieval adapter invoking image generation) without switching the client's active adapter. Invocable skills are explicitly allowlisted per adapter (<code>available_skills</code>); the calling adapter's LLM/retrieval pipeline is bypassed for that single turn and the skill's output is returned instead.	<code>docs/adapters/skills.md</code> , <code>server/services/chat_handlers/request_context_builder.py</code>
Automatic skill/intent routing	Optional per-adapter setting that infers which skill to invoke from plain natural language (no explicit <code>skill</code> field required), using an embedding pre-filter over each skill's example phrases followed by a constrained LLM confirmation call. Falls back to a normal conversational turn on any detection error. Disabled by default; opt-in per adapter.	<code>docs/adapters/skills.md</code> (Automatic Intent Detection), <code>config/config.yaml (skill_routing)</code>
Autocomplete / query suggestions	Real-time input suggestions drawn from intent-adapter <code>nl_examples</code> and skill <code>routing_examples</code> , ranked with fuzzy string matching (Levenshtein, Jaro-Winkler) and served from the same shared cache backend as the rest of the platform (SQLite/Redis/Memcached).	<code>docs/autocomplete-architecture.md</code> ,

Capability	Details	Source
		<code>config/config.yaml</code> (<code>autocomplete</code>)
Dynamic conversation-history token budgeting	Per-conversation history length is computed automatically from the active model's context window (read from provider-specific config, e.g. <code>num_ctx / n_ctx / context_window</code>) minus a reserved allowance for the system prompt and the model's own output-token setting, rather than a fixed message count. Adapters and per-request model selection can each apply their own budget. Old messages are automatically pruned once a session exceeds 120% of budget, with a warning surfaced to the client at 90%.	<code>docs/conversation_history.md</code>

2.10 Reference Client: OrbitChat

ORBIT ships a standalone chat client, distributed separately as an npm package, that connects to any ORBIT deployment. It is not part of the ORBIT server itself.

Capability	Details	Source
Distribution	Installable via <code>npm install -g orbitchat</code> ; runs as a CLI (<code>orbitchat</code>) or as a daemon (<code>orbitchat-daemon</code>)	<code>clients/orbitchat/README.md</code>
Feature set	Streaming responses, file upload, conversation threading, voice input/output, autocomplete, feedback capture, optional Auth0-based login	<code>clients/orbitchat/README.md</code>
Key-handling model	The browser never receives real API keys. Requests carry an <code>X-Adapter-Name</code> header; a bundled Express proxy resolves that name to the corresponding backend API key server-side before forwarding the request.	<code>clients/orbitchat/README.md</code> (Architecture, Security)
API-only mode	The same proxy layer can run without serving the bundled UI (<code>--api-only</code>), exposing a documented REST contract (<code>GET /api/adapters</code> , adapter-scoped chat/file/autocomplete endpoints) for a custom frontend to build against	<code>clients/orbitchat/README.md</code> (API-Only Mode)
Configuration	Per-deployment branding, adapter list, and feature toggles via a single <code>orbitchat.yaml</code> ; adapter secrets supplied separately via environment variables	<code>clients/orbitchat/README.md</code> (Configuring Adapters)

This positions OrbitChat as an optional, separately maintained reference client rather than a mandatory component — an organization can use ORBIT's API directly, build a custom frontend against the same proxy contract, or adopt OrbitChat as-is.

2.11 Test Coverage & Performance Testing

Capability	Details	Source
Automated test suite	214 test files and approximately 3,867 test functions under <code>server/tests/</code> , organized by subsystem (adapters, admin, auth, cache, CLI, config, datasources, embeddings, fault tolerance, inference, messaging, middleware, pipeline steps, retrievers, routes, services, threads, vector stores, vision, image, document generation, chat handlers, secrets, sound). Counted directly from the repository (<code>grep -rE "^\s*(async)?def test_"</code>); excludes <code>__pycache__</code> .	<code>server/tests/</code>
Load testing	Locust-based load testing (<code>locustfile.py</code>) with three user profiles (general, health-check, chat-focused), plus a shell-script runner (<code>run_performance_tests.sh</code>) covering basic, stress, health, chat, admin, and endurance scenarios	<code>server/tests/perf/README.md</code> , <code>server/tests/perf/locustfile.py</code>
Custom load/scenario testing	<code>advanced_performance_test.py</code> supports mixed, burst, and ramp load patterns with CSV/HTML reporting; <code>multi_user_load_test.py</code> simulates concurrent multi-tenant traffic across different adapters/API keys with per-adapter latency (p50/p95/p99) and success-rate breakdown	<code>server/tests/perf/advanced_performance_test.py</code> , <code>server/tests/perf/multi_user_load_test.py</code>
Rate-limit/throttle verification	<code>rate_limit_simulation.py</code> drives targeted traffic patterns (burst, sustained, random, throttle, quota-exhaustion) against the rate-limiting and throttling subsystem and reports observed behavior against configured limits	<code>server/tests/perf/rate_limit_simulation.py</code>
Memory-leak profiling	<code>memray_leak_test.py</code> runs a load scenario against the server under memray instrumentation and produces flamegraph/table reports of allocations, including those still live at process exit	<code>server/tests/perf/memray_leak_test.py</code>

The presence of dedicated fault-tolerance tests, load/stress tooling, rate-limit simulation, and memory-leak profiling indicates that performance and stability testing is treated as a first-class part of the test suite, not limited to functional unit tests.

2.12 Administration & Management UI/CLI


ORBIT includes a web-based admin panel and a companion command-line tool (both part of the server distribution, not separate products) for day-to-day management, gated by the RBAC permissions described in §2.6.

Capability	Details	Source
No-code adapter creation	Adapters can be created, edited, exported, and imported from the admin panel via a form-based flow, validated, and hot-reloaded without a server restart (including propagation across multi-worker deployments)	CHANGELOG.md (v2.14.0, v2.15.2, v2.15.7)
MCP server management	External MCP servers can be added, edited, and removed from the admin panel, with live reachability/tool diagnostics and hot reload — no config-file editing or restart required	CHANGELOG.md (v2.14.0–v2.15.7), docs/mcp_protocol.md
Admin CLI (<code>orbit</code>)	A command-line tool covering server lifecycle (start/stop/restart/status), authentication and user management, API key lifecycle (create/list/test/rename/deactivate/delete), system prompt management, CLI configuration, and admin operations (hot-reloading adapters or intent templates without a restart) — giving the same administrative surface as the web panel for scripting, automation, and CI/CD use	bin/orbit.py , bin/orbit.sh
Adapter SDK CLI (<code>adapter-sdk</code>)	A separate CLI for generating adapter YAML from the same spec registry the admin panel's adapter-creation form uses, including a machine-readable spec listing (<code>--list --json</code>) for tooling integration	bin/adapter-sdk.sh
Cost & audit dashboards	API key, adapter, provider, and call-type (chat/embedding/rerank/media) cost breakdowns with drill-down filtering, alongside the audit log viewer described in §2.6	CHANGELOG.md (v2.15.8–v2.15.9), docs/token-usage-and-cost-tracking.md

2.13 Onboarding & Documentation

Resource	Details	Source
Guided setup	An interactive installer (<code>install/setup.sh</code> , <code>install/wizard.py</code>) with configurable profiles, plus a prebuilt Docker path (<code>docker-compose.yml</code> , published images for Ollama/OpenAI/Gemini configurations) that starts a running instance with no manual configuration	install/ , docker/README.md
Step-by-step tutorial series	A <code>docs/tutorial.md</code> entry point linking a sequenced set of guides — first chat, chat with files, SQL/MongoDB/DuckDB querying, MCP tool calling, auto skill routing, creating API keys, admin panel tour, and troubleshooting	docs/tutorial.md , docs/tutorial/
Full documentation index	A categorized index (<code>docs/README.md</code>) covering architecture, adapters, data sources, security, and advanced/protocol topics, in addition to this capability matrix	docs/README.md
Runnable examples	Client code samples for the OpenAI-compatible API (Python/Node.js, including streaming) and a sample MCP server for hands-on testing	examples/

3. Architecture



Clients connect through the ORBIT API over REST, an OpenAI-compatible interface, MCP, A2A, or async MQ transports. Requests pass through auth/quota checks, then an adapter router that dispatches to chat/voice, document, structured-data, API, tool-agent, or generation handlers, which in turn call local or cloud model providers.

4. Platform-Specific Notes

4.1 LiteLLM

LiteLLM's core function is LLM API proxying — request/response translation, load balancing, and spend tracking across providers. It does not include built-in connectors to SQL/vector/NoSQL data sources; retrieval is typically implemented in a separate pipeline outside LiteLLM itself.

LiteLLM is distributed under an open-core model. Per LiteLLM's published pricing page (litellm.ai/pricing), the following are listed under a separately sales-quoted "Enterprise" tier rather than the open-source tier: SSO + SCIM, OIDC/JWT authentication, audit logs, secrets-manager integration with key rotation, org/team administration, a multi-region control plane, 24/7 support with SLAs, and self-hosted/air-gapped deployment. Evaluators comparing total cost of ownership should confirm current tier boundaries directly with LiteLLM, as pricing pages are subject to change.

ORBIT includes equivalents of the items listed above — OIDC/SSO, audit logging, secrets-manager integration, RBAC-based access administration, and self-hosted/air-gapped deployment — in its single Apache 2.0 distribution.

4.2 Open WebUI

Open WebUI is a chat frontend. It connects to Ollama natively and to any OpenAI-compatible endpoint (which can include self-hosted llama.cpp, vLLM, or LM Studio servers exposing that API). It does not include native connectors to SQL, vector, or NoSQL databases as retrieval sources; document RAG is limited to uploaded-file vector search.

For organizations evaluating a chat frontend specifically, ORBIT's comparable offering is OrbitChat (\$2.10) — a separately distributed reference client rather than a component of the ORBIT server itself. OrbitChat and Open WebUI are both frontends that can be pointed at the same class of backend model APIs; the difference in this comparison is in what sits behind them — ORBIT's data-source and governance layer versus a backend the deploying organization assembles separately.

ORBIT is a backend gateway rather than a chat frontend: it provides connection pooling for SQL databases, RabbitMQ-based async messaging, and multi-worker process scaling, none of which are part of Open WebUI's scope.

Regarding licensing: per Open WebUI's own documentation (docs.openwebui.com/license), versions through 0.6.5 are BSD-3-Clause (fully permissive). From v0.6.6 onward, the license adds a branding-protection clause requiring "Open WebUI" branding to remain visible unless the deployment has 50 or fewer users in a 30-day period, the user is a substantive contributor with written permission, or the user holds an enterprise license granting branding rights. Open WebUI's own documentation states this version of the license is not OSI-approved as "open source" due to that clause; source code remains publicly available regardless. A separate commercial/enterprise license is offered for white-labeling. Evaluators should confirm the license version in use, since organizations forking from v0.6.5 or earlier are not subject to the branding clause.

4.3 Commercial Managed Platforms (AWS Bedrock, Azure AI, Google Vertex AI)

These platforms provide managed model hosting and, in most cases, managed retrieval/vector-search services (e.g., AWS OpenSearch, Azure AI Search) tied to the vendor's own cloud infrastructure. They require outbound connectivity to the vendor's services and generally price on a per-token or per-request basis.

ORBIT can be deployed on the same cloud infrastructure or fully offline, and connects to databases the organization already operates without requiring data migration into a vendor-managed store. This is a deployment-model difference rather than a feature-completeness comparison — commercial platforms offer operational guarantees (SLAs, managed scaling, vendor support) that a self-operated deployment does not provide by default.