

ORBIT Capability Matrix

Overview

ORBIT (Open Retrieval-Based Inference Toolkit) is an open-source AI gateway, retrieval-augmented generation (RAG) engine, and agent-protocol host. It can be deployed on-premises, in a private cloud, or fully offline.


Models, data sources, adapters, and security settings are defined as YAML files under `config/`, rather than in server code. This allows configuration to be version-controlled and reviewed in a source repository like any other code artifact, diffed between commits, and promoted across development, staging, and production environments through the organization's existing CI/CD and change-management process — supporting typical DevSecOps practices for multi-environment deployment.

This document describes ORBIT's capabilities as implemented in its configuration files (`config/*.yaml`) and documentation (`docs/`), and compares them against three reference categories:

- **LiteLLM** — an LLM API proxy and router
- **Open WebUI** — a chat user interface for local/cloud LLMs
- **Commercial managed AI platforms** (e.g., AWS Bedrock, Azure AI, Google Vertex AI, IBM watsonx)

Each row in the comparison tables is intended to describe what each platform does, not to rank them. Where a claim about a comparison platform could not be verified against that platform's own documentation or pricing page, it has been noted as such or omitted.

1. Architecture



Clients connect through the ORBIT API over REST, an OpenAI-compatible interface, MCP, A2A, or async MQ transports. Requests pass through auth/quota checks, then an adapter router that dispatches to chat/voice, document, structured-data, API, tool-agent, or generation handlers, which in turn call local or cloud model providers.

2. Comparison Matrix

Dimension	ORBIT	LiteLLM	Open WebUI	Commercial Managed Platforms (Bedrock/Azure/Vertex)
Primary function	AI gateway + multi-source RAG + agent protocol host	LLM API proxy and router	Chat user interface for local/cloud LLMs	Managed cloud AI service
Inference backends	41 provider/runtime configurations (cloud APIs, local GGUF, vLLM, TensorRT-LLM, BitNet, AirLLM, TEE) — see note in §3.1 on default-disabled backends	Broad set of cloud/proxy API routes	Ollama-native; also reaches any OpenAI-compatible endpoint (OpenAI, or a self-hosted llama.cpp/vLLM/LM Studio server)	Vendor-operated model catalog
Data source integration	SQL (9 dialects), vector DBs (6 integrated), NoSQL, REST/GraphQL, web search (provider-native for Gemini/OpenAI/xAI, plus 7 external search backends usable with any LLM)	No built-in data-source connectors; typically paired with an external RAG pipeline	13 supported vector databases, web search with source citations, and hybrid (BM25 + vector) search with cross-encoder reranking; no built-in SQL/NoSQL database connectors as retrieval sources	Cloud-native connectors (S3/Blob/GCS) and managed vector search services
Document & media processing	Three document-parsing engines (Docling, MarkItDown, LLM/vision OCR); image, video, STT/TTS generation	Not included	8 supported document-extraction engines (including Tika, Docling, Mistral OCR); image generation (GPT-Image, Gemini, ComfyUI) and voice/audio (speech-to-text, text-to-speech)	Cloud OCR/vision services (e.g., Texttract, Form Recognizer)
Agent protocols	MCP (server and client host), Google A2A, RabbitMQ-based async messaging	Function/tool-call proxying	MCP support, OpenAPI-server tools, and a Python tools/functions framework; no documented A2A support	Vendor-specific agent frameworks
Conversational UX features	Cross-adapter skill invocation with optional automatic intent routing, query autocomplete from adapter/skill examples, and per-model dynamic conversation-history budgeting (see §3.9)	Not in scope — LiteLLM proxies inference calls and does not manage conversation state or UX	Model presets, cross-conversation memory of remembered facts, and Markdown-defined "Skills"; these are chat-UI/prompt-level features rather than ORBIT's server-side adapter-routing and dynamic per-model history budgeting	Varies by vendor SDK/console; not a standard cross-platform feature
Response feedback capture	Built-in per-message thumbs up/down with optional comment, stored per adapter/user, with a dedicated admin-panel analytics view	Not in scope — LiteLLM does not track end-user response ratings	Open WebUI has a per-message rating feature; whether it includes adapter-level trend/satisfaction analytics comparable to	Varies by vendor console; typically not exposed as a queryable per-response feedback dataset

Dimension	ORBIT	LiteLLM	Open WebUI	Commercial Managed Platforms (<i>Bedrock/Azure/Vertex</i>)
	(satisfaction trend, per-adapter ranking, negative-feedback triage) — see §3.9		§3.9 was not verified against its documentation for this comparison	
Security & identity	6-role RBAC (11 permissions), OIDC/SSO (Entra ID, Auth0), OS keyring integration, AES-256 file encryption — included in the open-source distribution	Proxy-level auth in the open-source tier; SSO/SCIM, OIDC/JWT auth, and audit logs are listed as part of the paid Enterprise tier per LiteLLM's published pricing (see §4.1)	RBAC with roles, groups, and per-resource permissions; SSO/OIDC/LDAP and SCIM 2.0; API keys — per its documentation, listed as part of the open-source product	Cloud IAM (AWS IAM, Azure RBAC, etc.)
Per-key usage limits	Each API key can be bound to one adapter and given its own daily/monthly quota and priority-based throttling, managed from the admin panel with live usage and reset controls (see §3.6)	Per-key rate/spend limits are also included in LiteLLM's open-source tier, using its own proxy-level design	Not a documented feature of the open-source product	Varies by vendor console; typically account- or project-level rather than per-key
Inference cost tracking	A local, configurable pricing table estimates cost on every request (text, media, and realtime-voice), with per-record and aggregated views (by model/provider/adapter/user/API key) in the admin panel — see §3.8	Also cost-tracked in the open-source tier, via LiteLLM's own maintained model-pricing map, with spend reporting by key/user/team	Included in the open-source product: an Analytics dashboard tracks token usage per message, model, and user. Per its own documentation, it does not calculate a dollar cost automatically — the admin multiplies the tracked token counts by their provider's price manually	Native, invoice-accurate billing through the vendor's own cloud billing console
Safety & moderation	Configurable moderation backends (OpenAI, Anthropic, Llama-Guard3, Shieldstral), plus two PII detectors: a local model (<code>privacy-filter</code>) and a Presidio analyzer integration	Integrates with external moderation/PII services (e.g., Presidio)	Basic moderation options	Managed cloud guardrail services
Hardware acceleration	CUDA, MPS, vLLM, TensorRT-LLM (FP8/INT8), SGLang, BitNet	Routes to external inference servers; does not run hardware-accelerated inference itself	Dependent on the Ollama backend it connects to	Managed cloud GPU infrastructure
Deployment model	Docker, Kubernetes, bare-metal, air-gapped, async worker processes	Docker, Kubernetes, serverless	Docker/Compose, Kubernetes/Helm, pip install, with Redis-backed horizontal scaling across multiple workers/nodes	Managed cloud SaaS/PaaS

Dimension	ORBIT	LiteLLM	Open WebUI	Commercial Managed Platforms (<i>Bedrock/Azure/Vertex</i>)
License	Apache 2.0 (single tier, no gated features)	Open-source core under a permissive license, with a separate paid Enterprise tier for the features listed in §4.1 (per litellm.ai/pricing)	BSD-3-Clause through v0.6.5; v0.6.6+ adds a branding-protection clause (not OSI-approved as "open source"), with a separate enterprise license for white-labeling — see §4.2	Proprietary cloud service

3. Capability Detail

3.1 LLM Gateway & Inference

ORBIT defines **41 inference provider/runtime configurations** in `config/inference.yaml`. Most are disabled by default and require an API credential or a locally running runtime to be usable — the count reflects configured integrations, not backends active in a given deployment.

Capability	Details	Source
Supported backends	Major cloud providers (OpenAI, Anthropic, Gemini, AWS Bedrock, Azure OpenAI, Vertex AI, Cohere, Mistral, IBM watsonx, and more), plus local/self-hosted runtimes (Ollama, vLLM, llama.cpp, and others)	<code>config/inference.yaml</code>
Local hardware acceleration	Support for running models efficiently on local GPUs, including optimized runtimes for high-throughput and low-bit/quantized models	<code>config/inference.yaml</code>
Confidential computing	Integration with a trusted-execution-environment (TEE) provider for confidential model inference	<code>config/inference.yaml</code>
Fault tolerance	Automatic circuit-breaking on repeated provider failures, with configurable retry timing and fallback to an alternate model	<code>config/config.yaml</code> (<code>fault_tolerance</code>)
Per-adapter provider overrides	Each adapter can be configured to use a different inference, embedding, or reranking provider and model than the system default	<code>config/adapters.yaml</code>

3.2 Data Sources & Retrieval

Capability	Details	Source
SQL databases	PostgreSQL, MySQL, MariaDB, SQLite, Supabase, Oracle, SQL Server, DuckDB, AWS Athena	<code>config/datasources.yaml</code>
Vector databases	Chroma, Qdrant, Milvus, Pinecone, Elasticsearch, Redis Vector	<code>config/datasources.yaml</code> , <code>config/stores.yaml</code>
NoSQL databases	MongoDB (Atlas or self-hosted), Cassandra	<code>config/datasources.yaml</code>
Web/API sources	Intent-based query generation against REST/JSON and GraphQL APIs (e.g. natural language mapped to calls against a REST endpoint or a GraphQL API), plus a web-scraping adapter (Firecrawl) with automatic content chunking for large pages	<code>config/adapters/intent.yaml</code>
Web search — provider-native	For supported LLM providers (Gemini, OpenAI, xAI), search and answer synthesis happen in a single call using the provider's own built-in search tool	<code>config/adapters/web-search.yaml</code> , <code>docs/adapters/web-search.md</code>
Web search — external providers	Any configured LLM can be paired with an external search engine (e.g., DuckDuckGo, Brave, Tavily, Google), so the search source and the answering model can be chosen independently	<code>config/adapters/web-search-providers.yaml</code> , <code>docs/adapters/web-search.md</code>
Embedding providers	OpenAI, Cohere, Mistral, Gemini, Voyage AI, and other embedding providers, plus local/self-hosted options	<code>config/embeddings.yaml</code>
Rerankers	Cohere, Jina, OpenAI, Anthropic, Voyage AI, and other reranking providers	<code>config/rerankers.yaml</code>
Hybrid retrieval scoring	Combines semantic similarity, keyword matching, and reranking to improve result relevance	<code>config/config.yaml</code> (<code>composite_retrieval</code>)
SQL query safety guard	Every generated SQL query is checked before it runs, to confirm it is a single, read-only, size-limited query — blocking any query that would modify data	<code>CHANGELOG.md</code> (v2.15.3)
Intent template validation	Query templates are checked for correctness when loaded, and can optionally require administrator approval before being used	<code>CHANGELOG.md</code> (v2.15.4)

Retrieval works through **intent-based query generation**: a natural-language question is matched to a pre-approved query template rather than having a general-purpose model generate a new database query on every request.

3.3 Document Processing

Capability	Details	Source
File formats	PDF, Word, PowerPoint, Excel, CSV, JSON, XML, HTML, plain text, EPUB, ZIP, images, and audio	config/config.yaml
Processing engines (in priority order)	Layout-aware document parsing, then document-to-markdown conversion, then AI vision/OCR as a fallback for harder documents	config/config.yaml (files.processing.processor_priority), config/ocr.yaml
Chunking strategies	Several methods for splitting documents into retrievable pieces (fixed-size, token-based, semantic, and structure-aware)	config/config.yaml (files.default_chunking_strategy)
Document security	File-type verification on upload and encryption of stored documents	config/config.yaml (files.processing.magika , files.encryption)
Storage backends	Local filesystem, AWS S3/MinIO, Azure Blob Storage, Google Cloud Storage	config/config.yaml (files.storage_backend)

3.4 Multimodal & Media Processing

Domain	Details	Source
Speech-to-text	Local, on-device transcription, plus cloud providers including OpenAI, Google, and Gemini	config/stt.yaml
Text-to-speech	Cloud providers (OpenAI, Google, Gemini, ElevenLabs) and local/self-hosted voice engines	config/tts.yaml
Audio content sanitization	Non-speech content (like code blocks or tables) is stripped out before text is converted to speech, so spoken responses stay natural	config/tts.yaml (tts.sanitize_content)
Image generation	OpenAI, Gemini, xAI, and other image-generation providers	config/image.yaml
Video generation	Gemini, xAI, and other video-generation providers	config/video.yaml
Vision/OCR analysis	OpenAI, Anthropic, Gemini, and other vision-capable providers, for reading text and content out of images	config/vision.yaml
Real-time speech-to-speech	Live, interruptible voice conversations (the user can talk over the assistant), rather than a separate transcribe-then-speak pipeline. Can be connected to ORBIT's own data sources so spoken answers are grounded in real information rather than the model's general knowledge — useful for live virtual-assistant scenarios.	config/adapters/qa.yaml , docs/adapters/grounded- realtime-voice.md

An administrator can restrict, per adapter, which specific models a user is allowed to select for image, video, audio, and web-search generation, rather than every adapter being locked to one fixed default. (`CHANGELOG.md` , v2.15.6)

3.5 Protocols

Protocol	Details	Source
REST / OpenAI-compatible	A standard chat API compatible with the widely-used OpenAI format, with streaming responses	<code>server/routes/</code>
Model Context Protocol (MCP)	Can act as an MCP server (exposing ORBIT's own capabilities to other tools) and as an MCP client (connecting to external MCP tools, such as a filesystem or GitHub)	<code>config/mcp_clients.yaml</code> , <code>docs/mcp_protocol.md</code>
MCP agentic tool-calling	The model can call one or more external tools, review the results, and reason over multiple steps within a single request — built directly into ORBIT, without relying on a separate agent-orchestration framework. This can be an explicit request feature, or enabled so a normal conversation automatically decides when a tool is needed. Supported across most major LLM providers.	<code>docs/adapters/mcp-agent.md</code>
Google A2A	Supports Google's Agent-to-Agent protocol for agent discovery and task execution between AI systems	<code>docs/a2a-protocol.md</code>
Async messaging	Message-queue-based request/response support for decoupled, asynchronous processing	<code>config/config.yaml (messaging)</code>
WebSockets	Real-time, two-way streaming for voice and live status updates	<code>/ws</code> , <code>/ws/metrics</code>
AI coding-agent integration	Because MCP and A2A are open, provider-neutral protocols, ORBIT can be connected to external agentic tools that speak either one. This is documented in detail for Claude (Desktop, Code, and SDK-based agents), which can call ORBIT as an MCP tool or address it as a peer agent over A2A. Other MCP- or A2A-compliant agent tools and coding assistants can connect the same way, though the level of support for each has not been individually verified for this comparison.	<code>docs/claude-agent-integration.md</code>

3.6 Security, Governance & Identity

Capability	Details	Source
Authentication	Industry-standard password hashing, secure session tokens, and integration with the operating system's credential store	docs/authentication.md
SSO	Single sign-on via Microsoft Entra ID and Auth0	config/config.yaml (auth.providers)
RBAC	Six built-in roles (e.g., admin, operator, auditor, analyst) covering eleven distinct permissions, so access can be scoped to what each role actually needs	docs/rbac-architecture.md
API key controls	Keys can be scoped to specific adapters, given daily/monthly usage quotas and rate limits, and restricted to specific users or email addresses	docs/api-keys.md
Per-key quotas & throttling	Each key's daily/monthly limit, remaining usage, and throttle priority can be viewed and edited from the admin panel, with one-click resets and a live usage report across all keys — no configuration-file editing required	server/routes/admin/api_keys.py , server/services/api_key_service.py , server/admin/admin_panel/tabs/api-keys.js
Identity blacklisting	Administrators can block users or accounts by pattern, immediately revoking their active sessions	config/config.yaml (auth.blacklist)
Audit logging	Records of user requests and administrative actions, stored in the organization's database of choice	config/config.yaml (internal_services.audit)
Secrets management	Credentials can be pulled from AWS Secrets Manager, Azure Key Vault, or GCP Secret Manager instead of being stored directly in configuration files	config/config.yaml (secrets_management)
Strict auth mode	An optional setting that requires every request to carry a verified user identity, not just a valid API key	config/config.yaml (auth.require_authenticated_user)

All items in this table are included in ORBIT's open-source (Apache 2.0) distribution; there is no separate paid tier.

3.7 Safety, Moderation & Privacy

Capability	Details	Source
Moderation backends	Configurable content-moderation models from OpenAI, Anthropic, and open-source options, to screen requests and responses	config/moderators.yaml , config/guardrails.yaml

Capability	Details	Source
Local PII detection	An on-premises model that detects personal information without sending any data outside the organization's network	<code>config/moderators.yaml (privacy_filter)</code>
PII detection via Presidio	Integration with Presidio, a self-hosted PII-detection tool supporting roughly 100 configurable data types (names, emails, etc.)	<code>config/moderators.yaml (presidio)</code> , <code>docs/security/pii-moderation.md</code>
Network security controls	Standard web security protections (e.g., cross-origin restrictions, content-security policy) plus rate limiting and throttling to prevent abuse	<code>config/config.yaml (security)</code>

3.8 Reliability & Performance

Capability	Details	Source
Multi-worker scaling	The server can run multiple worker processes to handle higher request volume	<code>config/config.yaml (performance)</code>
Caching	Configurable caching (SQLite, Redis, or Memcached) to speed up repeated requests	<code>config/config.yaml (internal_services.cache)</code>
Response optimization	Compressed responses and efficient re-validation to reduce bandwidth	<code>config/config.yaml (performance.compression , performance.etag_caching)</code>
Cost/token tracking	Tracks token usage and estimated dollar cost for every request, broken down by provider and model. Also prices non-text usage (image, video, TTS, STT, OCR) by its own discrete unit (per image, per second, per character, etc.), not just tokens.	<code>config/pricing.yaml , docs/token-usage-and-cost-tracking.md</code>
Pricing rate table & staleness tracking	Costs are estimated from a local, editable rate table covering every configured provider/model, rather than pulled from a provider billing API. The table records when it was last updated and flags itself as stale in the admin panel after a configurable number of days, so outdated rates don't silently look current. Distinguishes a model priced at genuinely \$0 (e.g. a local model) from one with no rate configured at all.	<code>config/pricing.yaml , server/services/pricing_service.py</code>
Cost dashboards	Per-request cost appears on individual audit records, and an aggregated Costs view breaks down total spend by model, provider, adapter, user, request type, and API key over a selectable time window	<code>docs/token-usage-and-cost-tracking.md</code>
Observability	Live telemetry, a metrics endpoint compatible with standard monitoring tools (Prometheus), and log rotation	<code>config/config.yaml</code>

3.9 Conversational Behavior & Adapter Intelligence

Capability	Details	Source
Skills (cross-adapter capability invocation)	One adapter can call on another adapter's capability mid-conversation — for example, a retrieval adapter triggering image generation — without the user switching adapters. Which skills an adapter can call is explicitly configured.	<code>docs/adapters/skills.md</code> , <code>server/services/chat_handlers/request_text_builder.py</code>
Automatic skill/intent routing	An optional setting that lets ORBIT infer which skill to invoke directly from a user's plain-language request, rather than requiring the request to name the skill explicitly. Off by default; enabled per adapter.	<code>docs/adapters/skills.md</code> (Automatic Intent Detection), <code>config/config.yaml</code> (<code>skill_routing</code>)
Autocomplete / query suggestions	Real-time input suggestions as a user types, drawn from example phrases configured for each adapter and skill	<code>docs/autocomplete-architecture.md</code> , <code>config/config.yaml</code> (<code>autocomplete</code>)
Dynamic conversation-history token budgeting	How much conversation history is kept is calculated automatically based on the model's context size, rather than a fixed number of messages — so older messages are trimmed only when needed, and the user is warned before that happens.	<code>docs/conversation_history.md</code>
Response feedback capture	Users can rate individual responses with a thumbs up/down and an optional comment. Ratings are tracked per adapter and per user, and shown in the admin panel as a satisfaction trend, a per-adapter ranking, and a list of recent negative feedback for follow-up — giving a direct signal for where responses need improvement.	<code>server/services/feedback_service.py</code> , <code>server/admin/admin_panel/tabs/feedback.js</code> , <code>docs/sqlite-schema.md</code> (<code>feedback</code> table)

3.10 Reference Client: OrbitChat

ORBIT ships a standalone chat client, distributed separately as an npm package, that connects to any ORBIT deployment. It is not part of the ORBIT server itself.

Capability	Details	Source
Distribution	Installable as a standard package, runnable as either a command-line tool or a background service	<code>clients/orbitchat/README.md</code>
Feature set	Streaming responses, file upload, conversation threading, voice input/output, autocomplete, feedback capture, optional login via Auth0	<code>clients/orbitchat/README.md</code>
Key-handling model	The end user's browser never sees the real backend API key — the request identifies which adapter to use, and a server-side component resolves that to the correct key before forwarding it	<code>clients/orbitchat/README.md</code> (Architecture, Security)

Capability	Details	Source
API-only mode	The same server-side component can run as a pure API layer without the bundled chat UI, for organizations building their own custom frontend	<code>clients/orbitchat/README.md</code> (API-Only Mode)
Configuration	Branding, available adapters, and feature toggles are set in a single configuration file per deployment	<code>clients/orbitchat/README.md</code> (Configuring Adapters)

This positions OrbitChat as an optional, separately maintained reference client rather than a mandatory component — an organization can use ORBIT's API directly, build a custom frontend against the same proxy contract, or adopt OrbitChat as-is.

3.11 Test Coverage & Performance Testing

Capability	Details	Source
Automated test suite	Roughly 3,900 automated tests across 214 test files, covering every major subsystem (adapters, authentication, data sources, inference, messaging, admin functions, and more)	<code>server/tests/</code>
Load testing	Simulated traffic tools that model realistic usage patterns, from everyday load to stress and endurance scenarios	<code>server/tests/perf/README.md</code> , <code>server/tests/perf/locustfile.py</code>
Custom load/scenario testing	Additional tooling for burst and ramping traffic patterns, and for simulating many concurrent tenants/API keys at once to measure latency and success rates per adapter	<code>server/tests/perf/advanced_performance_test.py</code> , <code>server/tests/perf/multi_user_load_test.py</code>
Rate-limit/throttle verification	Dedicated tooling that verifies the rate-limiting and throttling systems behave correctly under a variety of traffic patterns	<code>server/tests/perf/rate_limit_simulation.py</code>
Memory-leak profiling	A dedicated tool that runs load against the server while profiling memory allocations, to catch memory leaks before they reach production	<code>server/tests/perf/memray_leak_test.py</code>

The presence of dedicated fault-tolerance, load, rate-limit, and memory-leak testing tools indicates that performance and stability testing is treated as a first-class part of the test suite, not limited to basic functional checks.

3.12 Administration & Management UI/CLI

ORBIT includes a web-based admin panel and a companion command-line tool (both part of the server distribution, not separate products) for day-to-day management, gated by the RBAC permissions described in §3.6.

Capability	Details	Source
No-code adapter creation	New adapters can be created, edited, exported, and imported directly from the admin panel through a form — no manual configuration-file editing required — and changes apply immediately without restarting the server	CHANGELOG.md (v2.14.0, v2.15.2, v2.15.7)
MCP server management	External tool connections (MCP servers) can be added, tested, and removed from the admin panel, with live status checks — no config-file editing or restart required	CHANGELOG.md (v2.14.0–v2.15.7), docs/mcp_protocol.md
Admin CLI (orbit)	A command-line tool that covers the same administrative tasks as the web panel — starting/stopping the server, managing users and API keys, and reloading configuration — for scripting, automation, and CI/CD use	bin/orbit.py , bin/orbit.sh
Adapter SDK CLI (adapter-sdk)	A companion command-line tool for generating adapter configuration files, useful for scripted or bulk adapter setup	bin/adapter-sdk.sh
Cost & audit dashboards	Breakdowns of spending by API key, adapter, provider, and request type, alongside the audit log viewer described in §3.6	CHANGELOG.md (v2.15.8–v2.15.9), docs/token-usage-and-cost-tracking.md

3.13 Onboarding & Documentation

Resource	Details	Source
Guided setup	An interactive installer (install/setup.sh , install/wizard.py) with configurable profiles, plus a prebuilt Docker path (docker-compose.yml , published images for Ollama/OpenAI/Gemini configurations) that starts a running instance with no manual configuration	install/ , docker/README.md
Step-by-step tutorial series	A docs/tutorial.md entry point linking a sequenced set of guides — first chat, chat with files, SQL/MongoDB/DuckDB querying, MCP tool calling, auto skill routing, creating API keys, admin panel tour, and troubleshooting	docs/tutorial.md , docs/tutorial/
Full documentation index	A categorized index (docs/README.md) covering architecture, adapters, data sources, security, and advanced/protocol topics, in addition to this capability matrix	docs/README.md
Runnable examples	Client code samples for the OpenAI-compatible API (Python/Node.js, including streaming) and a sample MCP server for hands-on testing	examples/

4. Platform-Specific Notes

4.1 LiteLLM

LiteLLM's core function is LLM API proxying — request/response translation, load balancing, and spend tracking across providers. It does not include built-in connectors to SQL/vector/NoSQL data sources; retrieval is typically implemented in a separate pipeline outside LiteLLM itself.

LiteLLM is distributed under an open-core model. Per LiteLLM's published pricing page (litellm.ai/pricing), the following are listed under a separately sales-quoted "Enterprise" tier rather than the open-source tier: SSO + SCIM, OIDC/JWT authentication, audit logs, secrets-manager integration with key rotation, org/team administration, a multi-region control plane, 24/7 support with SLAs, and self-hosted/air-gapped deployment. Evaluators comparing total cost of ownership should confirm current tier boundaries directly with LiteLLM, as pricing pages are subject to change.

ORBIT includes equivalents of the items listed above — OIDC/SSO, audit logging, secrets-manager integration, RBAC-based access administration, and self-hosted/air-gapped deployment — in its single Apache 2.0 distribution.

4.2 Open WebUI

Open WebUI is a chat frontend. It connects to Ollama natively and to any OpenAI-compatible endpoint (which can include self-hosted llama.cpp, vLLM, or LM Studio servers exposing that API). Per its own documentation (docs.openwebui.com/features), its retrieval layer is substantial for a chat frontend — 13 supported vector databases, hybrid (keyword + vector) search with cross-encoder reranking, and multiple document-extraction engines — though it does not include native connectors to SQL or NoSQL databases as retrieval sources the way ORBIT's intent-based retrieval does.

For organizations evaluating a chat frontend specifically, ORBIT's comparable offering is OrbitChat (\$3.10) — a separately distributed reference client rather than a component of the ORBIT server itself. OrbitChat and Open WebUI are both frontends that can be pointed at the same class of backend model APIs; the difference in this comparison is in what sits behind them — ORBIT's data-

source and governance layer versus a backend the deploying organization assembles separately.

ORBIT is a backend gateway rather than a chat frontend: it provides connection pooling for SQL databases, RabbitMQ-based async messaging, and multi-worker process scaling, none of which are part of Open WebUI's scope.

Regarding cost visibility: per Open WebUI's own documentation (docs.openwebui.com/features/administration/analytics), its open-source Analytics dashboard tracks token usage per message, model, and user, but does not calculate a dollar cost automatically — it documents the formula for an administrator to multiply tracked token counts by their provider's price by hand. This differs from ORBIT's and LiteLLM's approach of resolving an estimated dollar cost automatically from a maintained pricing table.

Regarding licensing: per Open WebUI's own documentation (docs.openwebui.com/license), versions through 0.6.5 are BSD-3-Clause (fully permissive). From v0.6.6 onward, the license adds a branding-protection clause requiring "Open WebUI" branding to remain visible unless the deployment has 50 or fewer users in a 30-day period, the user is a substantive contributor with written permission, or the user holds an enterprise license granting branding rights. Open WebUI's own documentation states this version of the license is not OSI-approved as "open source" due to that clause; source code remains publicly available regardless. A separate commercial/enterprise license is offered for white-labeling. Evaluators should confirm the license version in use, since organizations forking from v0.6.5 or earlier are not subject to the branding clause.

4.3 Commercial Managed Platforms (AWS Bedrock, Azure AI, Google Vertex AI)

These platforms provide managed model hosting and, in most cases, managed retrieval/vector-search services (e.g., AWS OpenSearch, Azure AI Search) tied to the vendor's own cloud infrastructure. They require outbound connectivity to the vendor's services and generally price on a per-token or per-request basis.

ORBIT can be deployed on the same cloud infrastructure or fully offline, and connects to databases the organization already operates without requiring data migration into a vendor-managed store. This is a deployment-model difference rather than a feature-completeness comparison — commercial platforms offer operational guarantees (SLAs, managed scaling, vendor support) that a self-operated deployment does not provide by default.